


DPO round 1: clean held-out, n=8 per problem, $\beta=0.1$, ref anchored to B1++

Clean held-out, bare prompt: recipe progression



One DPO round on the SFT checkpoint

