


DPO-r1 / DPO-r2 implicit reward structure: the strongest gradient is *against cheating**, not *for compliance**



Within the chosen pool, compliant $\wedge$ pass is the priority-1 example; compliant $\wedge$ fail is the fallback. Within the rejected pool, cheating is priority-1; non-compliant $\wedge$ fail is the fallback. The two top-priority quadrants dominate the loss — "the win" and "cheating" — leaving compliant $\wedge$ fail and incompetent rule-break with much smaller signal. This is the asymmetry that likely drove the round-2 bailout behavior.