


Pareto frontier on LCB clean: pass vs compliance



Each point is one recipe evaluated on the same 17-problem clean held-out set, $n=8$ samples each = 136 attempts. AST-rechecked from saved sources; both axes normalized to /136 (gen-errors and truncations count as fails / non-compliant). Pareto frontier connects {base, R-SFT, R-SFT+, DPO-r1} — moving up-left along it trades pass rate for compliance. Vanilla SFT is dominated by R-SFT (less pass AND less compliance). DPO-r2 is dominated by DPO-r1 — round 2 strictly worsens both. Base sits in a dashed uncertainty box: the lower-left corner is the measured value; the box extends to the strict upper bound if all 24 unmeasured attempts (3 problems lost mid-resume) were maximally favorable on both axes — other adapters showed ~ 0 compliance there, so the true base position is near the lower corner.