


Same 66 demos, same SFT recipe — removing rationale prose nearly doubles cheating



Pre-DPO sampling on the shared 45-problem LCB pool. The rationale-SFT model produces fewer violations and far less cheating; the stripped model produces more violations and 92% more cheating before any DPO step runs.