


Same 66 demos, same SFT recipe — removing rationale prose nearly doubles cheating



Pre-DPO sampling on the shared 45-problem LCB pool. "Cheating" is the subset of non-compliant gens that ALSO passed tests. With the rationale prose stripped, the model produces 78% more rule-breaks and 92% more cheating before any DPO step runs.