



Vanilla SFT + DPO is worse than rationale-SFT + DPO — same DPO recipe, only the SFT prose differs



Both adapters trained with identical recipe (66 demos, $\beta=0.1$, $lr=5e-6$, 3 epochs, same 45-problem DPO pool). The only difference is whether the SFT targets included the rationale prose. Result: on the truly-held-out clean set, removing rationale prose costs 43 pts of compliance and quadruples cheating.